

Overview

Preparing Anaconda

Preparing R

Getting the data

Importing the data

Robust outlier filtering

Benchmark forecasting the time series data

TensorFlow LSTM - Importing keras into R

LSTM data preparation

Building the LSTM

Training the LSTM

Compute performance metrics

Plot the LSTM train observations and train predictions

LSTM loss functions

Plot all observations, LSTM train predictions and testing predictions

# LSTM time series forecasting with TensorFlow

Dennis W. Hallema

Disclaimer: Use at your own risk. No responsibility is assumed for a user's application of these materials and related materials. These materials and referenced materials do not constitute an endorsement of any kind, not of any software, data set, method nor of anything else.

Description: These materials include a demonstration of the LSTM (Long-Short Term Memory) deep learning time series forecasting model in Keras/TensorFlow.

References:

- Hallema, D. W. (2024). Interactive Workshop, Process-Agnostic Water Quality Modeling DRI-ERDC Technology Transfer. Desert Research Institute, Las Vegas, NV, June 20, 2024.
- Hallema, D. W. (2024). LSTM (Long Short-Term Memory) forecasting of reservoir inflow and water quality in critical water supply areas. USACE-ERDC Broad Agency Announcement, No. W912HZ2320025: Drivers and disturbances of reservoir water temperature and water quality, FY 24 Initiative 3: AI-based assessment of regional water quality patterns in USACE reservoirs across the western United States (Technical Report). Desert Research Institute, July 2024, Las Vegas, Nevada, 48 p.

# Overview

We can use time series forecasting to predict the volume of water coming into the reservoir from upstream. A combination of weather models and distributed hydrologic models is useful for time scales shorter than two weeks, and climate models help with seasonal and annual predictions. However, for biweekly to seasonal time scales, LSTM time series forecasting is useful because it requires little data. This has several benefits:

1. Streamflow prediction and flood forecasting: An accurate prediction of water levels in rivers and reservoirs is critical for advance flood warning. In addition, this serves fish flow management. Fish flow mapping assists resource managers in linking aquaculture, fish farming and conservation policies, and flow regulation is an important element of fish flow regulation.
2. Climate impact analysis: Sea Surface Temperatures (SSTs) are used in climate modeling, but can also be linked directly to flow or stream temperature series forecasting models. This is done by incorporating area-averaged time series of SSTs into the LSTM.
3. Potential evapotranspiration (PET) forecasting. Many distributed hydrologic models, like the Storm Water Management Model (SWMM) and HEC-HMS, and linear hydrologic reservoir models can benefit from accurate evaporation data for calibration. PET co-varies with the diurnal temperature cycle, and this knowledge can help generate a PET time series forecast.

**NOTE:** Forecasting refers to the estimation of future values for a time series based on historical data. Prediction refers to the estimation of unknown past or future values in general. We speak of prediction when we train-test a model on historical data.

## Preparing Anaconda

TensorFlow and Keras are both written in Python, so we will need a Python distribution like Anaconda.

### 1. Download and install Anaconda.

- Install for Just Me (recommended)
- **IMPORTANT:** Check the box Register Anaconda3 as my default Python 3.X
- Check the box Clear the package cache upon completion
- (Check the box Add Anaconda to path variable)
- (Wait several minutes for installation to complete)

### 2. Add Anaconda to the PATH variable. In Windows:

- Press *Windows* key > Type environment variables > Enter
- Click Environment Variables
- Under User Variables, click Path > Edit
- If you don't see Anaconda listed, click New and add a variable  
C:\Users\Name\Anaconda3\Scripts where Name is your Windows user profile name.

Next, to install the TensorFlow and keras packages in the Anaconda environment, proceed as follows.

### 1. Press the *Windows* key and type Anaconda Navigator

2. Launch the PowerShell Prompt . This opens a PowerShell session in the Anaconda base environment, indicated by the prompt (base) PS C:\Users\Name> .
3. At this point you can elect to create a new Anaconda environment if you like. In this guide we'll install keras and a full version of tensorflow in the Anaconda base environment:
  - (base) C:\Users\Name> pip install tensorflow
  - (base) C:\Users\Name> pip install keras
4. Let's test if Tensorflow and Keras were installed correctly:
  - Launch Python: (base) C:\Users\Name> python
  - Import the Tensorflow package: >>> import tensorflow as tf
  - Import the Keras package: >>> import keras If there are no errors but only some warnings, the installation was likely succesful.

## Preparing R

Install the following packages using the command `install.packages()` as shown below.

```
# (Uncomment the following lines and run to install the packages)
# install.packages("dplyr")
# install.packages("readr")
# install.packages("pracma")
# install.packages("reticulate") # Python interface
### install.packages("devtools")
```

We also need the R package `keras3` so we can use R as an interface for Keras. When prompted, select "1" to update all keras dependencies. After the installer finishes, check the package version to ensure it was installed correctly.

```
# install.packages("keras3")
# packageVersion("keras3")

# (Try the following if the Keras installation fails)
### install.packages("devtools")
### devtools::install_github("rstudio/keras")
```

## Getting the data

Copy `usgs_14210000_2014_2024.csv` to the same directory as this `.Rmd` file (the R Markdown Notebook).

(If you first need to download the data, open the following link and wait for the download to complete:

[https://nwis.waterdata.usgs.gov/nwis/uv?](https://nwis.waterdata.usgs.gov/nwis/uv?cb_all_=on&cb_00010=on&format=rdb&site_no=14210000&legacy=1&period=&begin_date=2014-04-01&end_date=2024-04-01)

`cb_all_=on&cb_00010=on&format=rdb&site_no=14210000&legacy=1&period=&begin_date=2014-04-01&end_date=2024-04-01` ([https://nwis.waterdata.usgs.gov/nwis/uv?](https://nwis.waterdata.usgs.gov/nwis/uv?cb_all_=on&cb_00010=on&format=rdb&site_no=14210000&legacy=1&period=&begin_date=2014-04-01&end_date=2024-04-01)

`cb_all_=on&cb_00010=on&format=rdb&site_no=14210000&legacy=1&period=&begin_date=2014-04-01&end_date=2024-04-01`) > Right-click > Save As `usgs_14210000_1984_2024.csv` )

The data file contains a tab-separated USGS time series for the stream gage

USGS 14210000 CLACKAMAS R. near ESTACADA Oregon covering the period 2014-4/2024-4. Variables include stream temperature, dissolved oxygen, and turbidity, among others. To view the location where these data

were collected, browse to Google Maps (or Bing Maps) and copy-paste the coordinates:  
45°18'00.0"N 122°21'10.0"W .

**Question:** What do you notice about the location of the stream gage?

Open the USGS data file in Excel. It looks like this:

```

# ----- WARNING -----
# Some of the data that you have obtained from this U.S. Geological Survey database
# may not have received Director's approval. Any such data values are qualified
# as provisional and are subject to revision. Provisional data are released on the
# condition that neither the USGS nor the United States Government may be held liable
# for any damages resulting from its use.
#
# Additional info: https://waterdata.usgs.gov/provisional-data-statement/
#
# File-format description: https://help.waterdata.usgs.gov/faq/about-tab-delimited-output
# Automated-retrieval info: https://help.waterdata.usgs.gov/faq/automated-retrievals
#
# Contact: gs-w_waterdata_support@usgs.gov
# retrieved: 2024-04-16 11:09:02 EDT (sdww01)
#
# Data for the following 1 site(s) are contained in this file
# USGS 14210000 CLACKAMAS RIVER AT ESTACADA, OR
# -----
#
# Data provided for site 14210000
#
# TS parameter statistic Description
# 114314 00060 00003 Discharge, cubic feet per second (Mean)
# 114315 00010 00001 Temperature, water, degrees Celsius (Maximum)
# 114316 00010 00002 Temperature, water, degrees Celsius (Minimum)
# 114317 00010 00003 Temperature, water, degrees Celsius (Mean)
# 114318 00095 00001 Specific conductance, water, unfiltered, microsi
emens per centimeter at 25 degrees Celsius (Maximum)
# 114319 00095 00002 Specific conductance, water, unfiltered, microsi
emens per centimeter at 25 degrees Celsius (Minimum)
# 114320 00095 00003 Specific conductance, water, unfiltered, microsi
emens per centimeter at 25 degrees Celsius (Mean)
# 114321 00300 00001 Dissolved oxygen, water, unfiltered, milligrams
per liter (Maximum)
# 114322 00300 00002 Dissolved oxygen, water, unfiltered, milligrams
per liter (Minimum)
# 114323 00300 00003 Dissolved oxygen, water, unfiltered, milligrams
per liter (Mean)
# 114324 00400 00001 pH, water, unfiltered, field, standard units (Ma
ximum)
# 114325 00400 00002 pH, water, unfiltered, field, standard units (Mi
nimum)
# 114326 00400 00008 pH, water, unfiltered, field, standard units (Me
dian)
# 114330 00045 00006 Precipitation, total, inches (Sum), Faraday Lake
precip
# 216058 63680 00008 Turbidity, water, unfiltered, monochrome near in
fra-red LED light, 780-900 nm, detection angle 90 +-2.5 degrees, formazin nephelometric un
its (FNU) (Median)
# 216059 63680 00001 Turbidity, water, unfiltered, monochrome near in
fra-red LED light, 780-900 nm, detection angle 90 +-2.5 degrees, formazin nephelometric un
its (FNU) (Maximum)
# 216060 63680 00002 Turbidity, water, unfiltered, monochrome near in

```



# Importing the data

As shown above, the data comes with multiple lines of comments preceded by a `#`, and a header line. We will use the `readLines` function to read in the whole file line-by-line, and then `read_table` (`readr` package), which allows us to automatically organize the raw data into columns.

To ensure R can find our data, copy the file `usgs_14210000_1984_2024.csv` in the same folder as the R Markdown notebook. Next, return to RStudio and select the `Session` menu > `Set working directory > To source file location`. If we attempt to open a file, R will now look in the directory where the source file is also located.

We will now import the data. Note that R may report some parsing failures, but that is not a problem.

```
# Load packages
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
##
##   filter, lag
```

```
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(readr)

# Read raw data file
file <- "usgs_14210000_1984_2024.csv"
cleanlines = readLines(file)

# Remove comments and first header line
for(i in c("#","$s")){ cleanlines <- sub(paste0(i, ".*"), "", cleanlines) }

# Automatically organize the raw data into columns
datraw <- read_table(cleanlines)
```

```
## Warning: 3657 parsing failures.
## row col   expected      actual      file
##   1  -- 49 columns 36 columns literal data
##   2  -- 49 columns 36 columns literal data
##   3  -- 49 columns 36 columns literal data
##   4  -- 49 columns 36 columns literal data
##   5  -- 49 columns 36 columns literal data
## ... ..
## See problems(...) for more details.
```

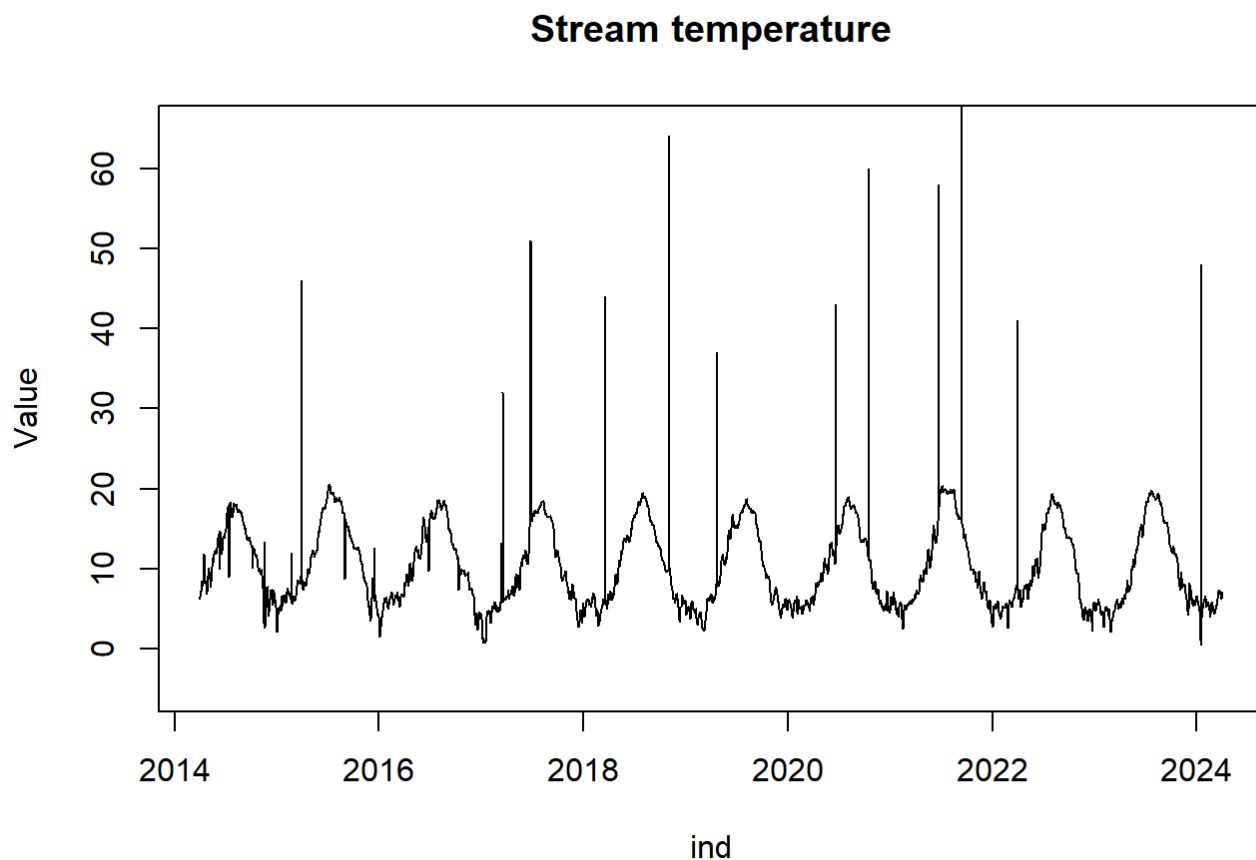
Copy the column with the label `Temperature, water, degrees Celsius (Mean)`. According to the comment lines at the beginning of the data file (the lines marked with `#`), we need the column with the index `114317_00010_00003`.

```
# Copy the column of interest
dat <- data.frame(ind = datraw$datetime, Value = as.numeric(datraw$`114317_00010_00003`))

# Inspect the first few lines
head(dat)
```

```
##           ind Value
## 1 2014-04-01   6.2
## 2 2014-04-02   6.4
## 3 2014-04-03   6.5
## 4 2014-04-04   6.8
## 5 2014-04-05   6.9
## 6 2014-04-06   7.2
```

```
# Visualize the data
plot(dat, main = "Stream temperature", type = "l", ylim = c(-5,65))
```



**Question:** What stands out in this plotted time series?



# Robust outlier filtering

Outliers can represent true values or result from error (observational error, data error, etc.) To separate between true extremes from extreme values due to error, we apply a Hampel filter. The Hampel filter computes the median value for a running window between  $-k$  and  $k$  time steps away of each point, and determines the median of absolute deviations (MAD) from this running median. If a data point lies more than  $t_0 \times \text{MAD}$  away from the running median, we may presume that its value is not a true value. We replace those with the running median. Let's identify the erroneous extreme values.

```
# Load the PRACMA package
library(pracma)

# non NA values index
ValInd <- which(!is.na(dat$Value))

# Use the Hampel filter (mean absolute deviation from median) to find the rows that have o
utliers
OutlierInd <- ValInd[pracma::hampel(dat$Value[ValInd],k=7, t0=7)$ind]
# dat$Value[OutlierInd]

# Number of outliers detected
nrow(dat[OutlierInd,])
```

```
## [1] 23
```

```
# View all outliers
dat[OutlierInd,]
```

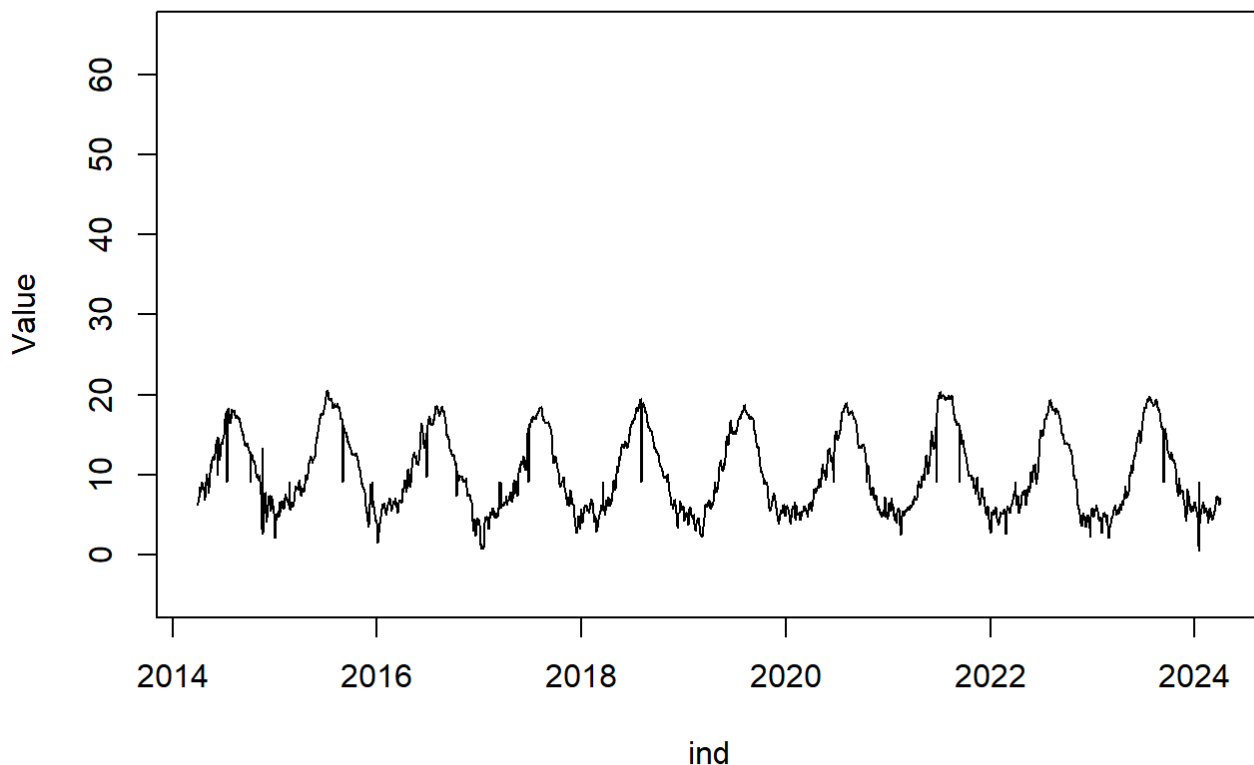
```
##           ind Value
## 17    2014-04-17  11.8
## 106   2014-07-15   9.0
## 190   2014-10-07  10.1
## 330   2015-02-24  12.0
## 365   2015-03-31  46.0
## 520   2015-09-02   8.7
## 625   2015-12-16  12.6
## 954   2016-11-09   9.3
## 1079  2017-03-14  13.2
## 1084  2017-03-19  32.0
## 1183  2017-06-26  51.0
## 1450  2018-03-20  44.0
## 1587  2018-08-04  18.8
## 1677  2018-11-02  64.0
## 1849  2019-04-23  37.0
## 2272  2020-06-19  43.0
## 2391  2020-10-16  60.0
## 2640  2021-06-22  58.0
## 2724  2021-09-14  71.0
## 2922  2022-03-31  41.0
## 3451  2023-09-11  15.8
## 3453  2023-09-13  15.8
## 3579  2024-01-17  48.0
```

Next, we remove the identified outliers and visualize the resulting time series.

```
# Replace outliers
for (i in 1:length(OutlierInd)) {
  dat$Value[OutlierInd] <- median(dat$Value, na.rm=T)
}

# Plot the time series again
plot(dat, main = "Stream temperature (MAD filtered)", type = "l", ylim = c(-5,65))
```

## Stream temperature (MAD filtered)



**Question:** What Hampel window size did we use here? And what filter threshold? Tip: Open the help documentation for the Hampel function using `?hampel` .

**Question:** What is the advantage of using a Hampel median absolute deviation (MAD) filter to remove outliers as opposed to filters based on a running mean?

## Benchmark forecasting the time series data

We will get to LSTM deep learning forecasting in a bit. To understand the advantage of LSTMs, it can help to look at other forecasting methods (benchmarking) first. The simplest benchmark is based on the assumption that nothing changes, i.e. we compute an expected value from the mean of all the previous values. This yields a rough estimate of what we can expect, and we can add a confidence interval, which we expect to increase as a function of time.

```
# Load the forecast package
library(forecast)
```

```
## Registered S3 method overwritten by 'quantmod':
```

```
##   method          from
```

```
##   as.zoo.data.frame zoo
```

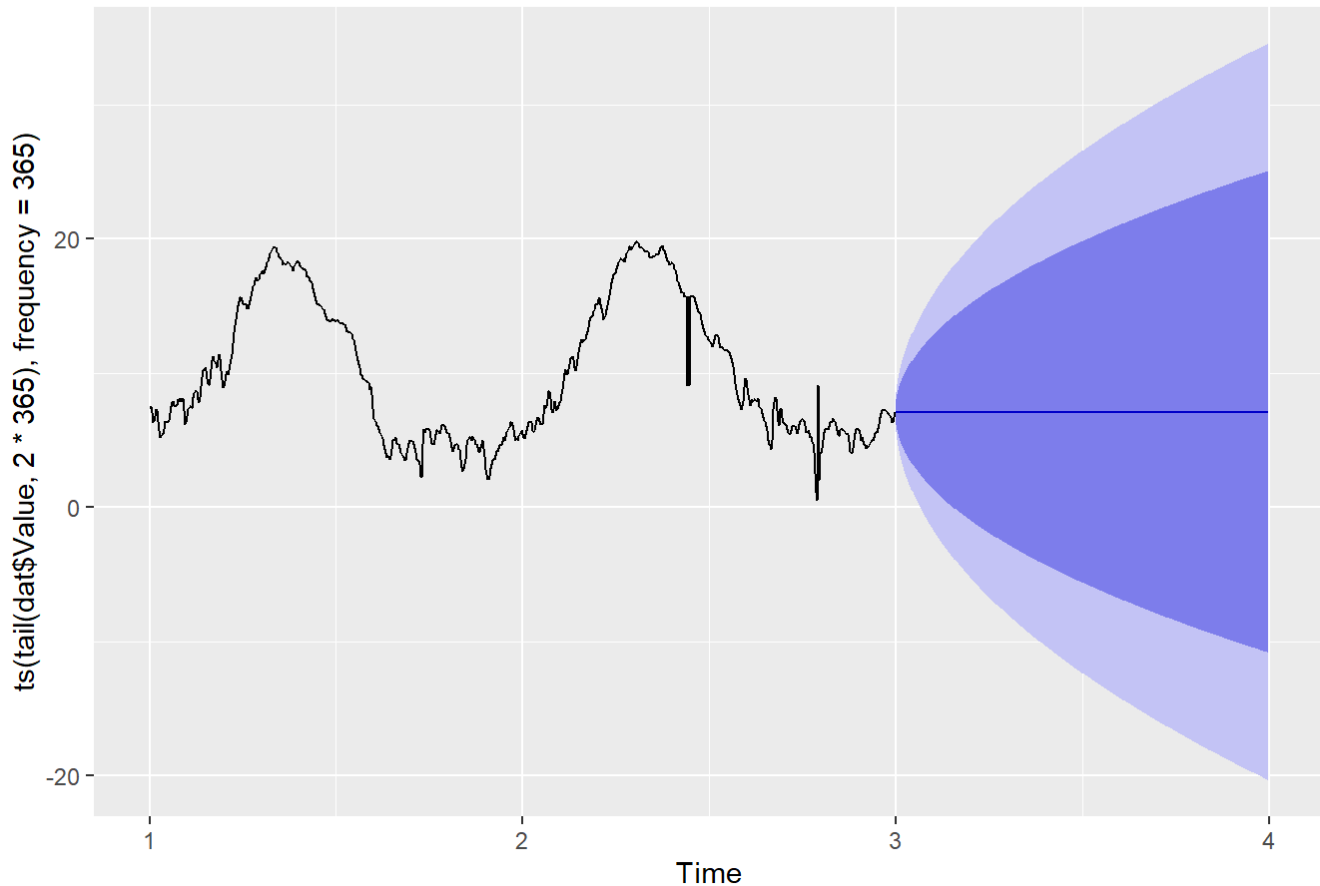
```
# Naive forecast
```

```
ndat <- naive(ts(tail(dat$Value, 2*365), frequency = 365), h = 365)
```

```
# Plot and summarize the forecasts
```

```
autoplot(ndat)
```

### Forecasts from Naive method



**Question:** What does the naive function do? Tip: use the command `?naive`.

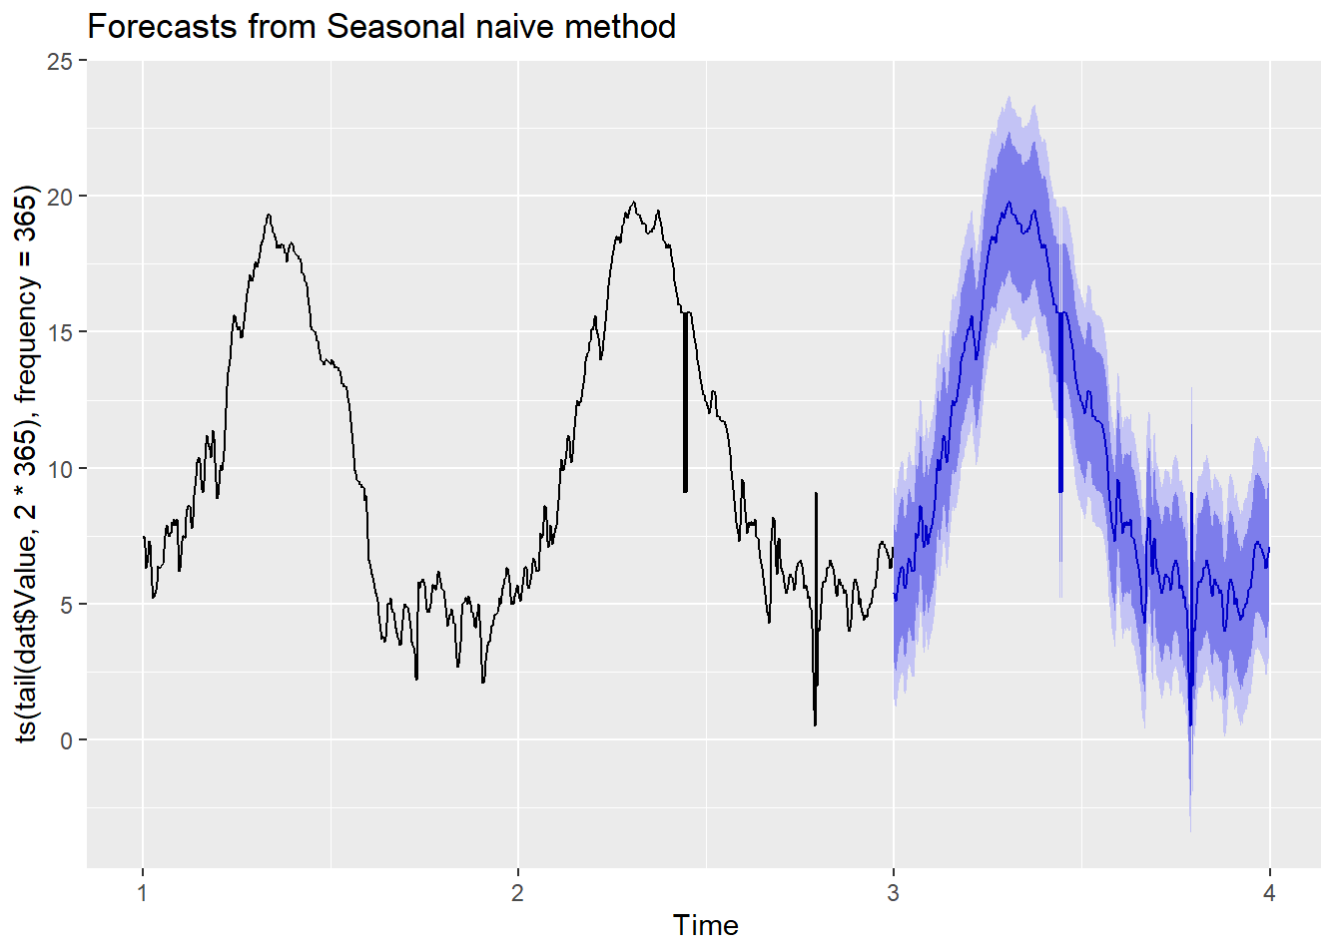
Now let's add a seasonal component to the forecasting model.

```
# Seasonal naive forecast
```

```
sndat <- snaive(ts(tail(dat$Value, 2*365), frequency = 365), h = 365)
```

```
# Plot and summarize the forecasts
```

```
autoplot(sndat)
```



**Question:** How does the seasonal naive streamflow temperature forecast method compare to the naive forecast method?

## ARIMA models

Naive models are Autoregressive Integrated Moving Average (ARIMA) models. ARIMA uses the lag and shift in time series data to predict future patterns. The model was first proposed in 1976 and has gained popularity in economic forecasting and short-term business performance forecasting. To understand how ARIMA works, we need to discuss its predecessors because ARIMA combines components of these earlier models.

- Autoregressive (AR) models: multiple regression using lagged observations as predictors.
- Moving Average (MA) models: multiple regression using lagged residuals (errors) as predictors.
- Autoregressive Moving Average (ARMA) models: multiple regression using both lagged observations and lagged residuals as predictors. Accepts only stationary data.
- Autoregressive Integrated Moving Average (ARIMA{p,d,q}) models: This is generalization of the ARMA model with differencing. This allows ARIMA to account for non-stationarity in the data. The ARIMA parameters (p,d,q) are non-negative integers describing the autoregressive, differencing, and moving average terms:
  - p: order of the autoregressive (AR) model, i.e. the number of lagged observations
  - d: degree of differencing (I), i.e. the number of differenced observations
  - q: order of the moving-average (MA) (or seasonal) model In addition to (p,d,q), ARIMA models can also include intercepts or constants (c), in some cases referred to as the “drift”.

These are some common and less common ARIMA model configurations:

- ARIMA(0,0,0): white noise model
  - ARIMA(0,0,1): first-order moving average (MA) model. Same as MA(1)
  - ARIMA(0,1,0): random walk (if this model uses a constant  $c$  it describes a random walk with drift). Same as  $I(1)$ . This is essentially a stochastic process.
  - ARIMA(0,1,1) without constant: simple exponential smoothing
  - ARIMA(0,1,1) with constant: simple exponential smoothing with growth
  - ARIMA(0,1,2): Damped Holt's model
  - ARIMA(0,2,1) or (0,2,2) without constant: linear exponential smoothing
  - ARIMA(0,2,2): double exponential smoothing
  - ARIMA(0,3,3): triple exponential smoothing
  - ARIMA(1,0,0): first-order autoregressive model. If the series is stationary and autocorrelated, and can be predicted as a multiple of the last value (plus constant). Same as AR(1)
  - ARIMA(1,1,0): differenced first-order autoregressive model. Used if the errors of the random walk model are autocorrelated but without seasonality.
  - ARIMA(1,1,2) without constant: damped-trend linear exponential smoothing
  - ARIMA(2,1,1): second order autoregressive, first order moving average, differenced once
- auto.arima(): The auto.arima() function in the forecast package for R can automatically select the best ARIMA model. It conducts a search over possible models within a set of constraints, and returns the best ARIMA model according to the AICC (or AIC or BIC) value.

**Question:** What type of ARIMA(p,d,q) configuration describes a naive model? And a seasonal naive model? Tip: ?naive

## TensorFlow LSTM - Importing keras into R

Now we arrive at the deep learning model. First, we will setup R so we can use it as a Python interface.

```
# Set the Python path to Anaconda Python
Sys.setenv(RETICULATE_PYTHON = paste0("C:/Users/", Sys.getenv("USERNAME"), "/AppData/Local/anaconda3/python.exe"))
Sys.setenv(PYTHON = paste0("C:/Users/", Sys.getenv("USERNAME"), "/AppData/Local/anaconda3/python.exe"))
library(reticulate)
reticulate::py_config()
```

```
## python:      C:/Users/dhallema/AppData/Local/anaconda3/python.exe
## libpython:   C:/Users/dhallema/AppData/Local/anaconda3/python312.dll
## pythonhome:  C:/Users/dhallema/AppData/Local/anaconda3
## version:     3.12.4 | packaged by Anaconda, Inc. | (main, Jun 18 2024, 15:03:56) [MS
C v.1929 64 bit (AMD64)]
## Architecture: 64bit
## numpy:       C:/Users/dhallema/AppData/Local/anaconda3/Lib/site-packages/numpy
## numpy_version: 1.26.4
##
## NOTE: Python version was forced by RETICULATE_PYTHON
```

```
# Import keras into Python
reticulate::import("keras")
```

```
## Module(keras)
```

```
library(keras3)
```

## LSTM data preparation

Next, let's collect the values of interest and impute the NAs. Unlike other models, neural networks are designed to demote any data points without predictive power or relation to adjacent data points. In data sets like this where 0 is not a common value in the data, the best practice is to impute all NAs with 0 .

```
# Load the reticulate and keras packages
library(reticulate)
library(keras3)

# Generate model input
ts <- data.frame(
  date = as.character(dat$ind),
  val = as.numeric(dat$value)
)

# Impute NAs
ts$val[is.na(ts$val)] <- 0

# Group by date
ts_grouped <- ts %>% group_by(date) %>%
  summarise(val=mean(val)) %>%
  as.data.frame()
dt <- ts_grouped

# Center and scale data
values <- data.frame(scale(dt$val))
```

In the last step above, we also center and scale the data. To train the LSTM, we will need to select the data we will use for training. We'll set a small portion of the data aside so we can test the performance of the model.

```

# Function to Lag and split data into train and test
lagsplit_data <- function(tseries, test_size = 0.2, lookback = 30) {
  data_raw <- as.matrix(tseries) # convert to matrix
  data <- array(dim = c(0, lookback, ncol(data_raw)))

  # Lag the data lookback times and arrange into columns
  for (i in 1:(nrow(data_raw) - lookback)) {
    data <- rbind(data, data_raw[i:(i + lookback - 1), ])
  }

  test_set_size <- round(test_size * nrow(data))
  train_set_size <- nrow(data) - test_set_size

  x_train <- data[1:train_set_size, 1:(lookback - 1), drop = FALSE]
  y_train <- data[1:train_set_size, lookback, drop = FALSE]

  x_test <- data[(train_set_size + 1):nrow(data), 1:(lookback - 1), drop = FALSE]
  y_test <- data[(train_set_size + 1):nrow(data), lookback, drop = FALSE]

  cat(paste('x_train.shape = ', dim(x_train), '\n'))
  cat(paste('y_train.shape = ', dim(y_train), '\n'))
  cat(paste('x_test.shape = ', dim(x_test), '\n'))
  cat(paste('y_test.shape = ', dim(y_test), '\n'))

  # Return a list with train and test datasets lagged for each lookback period
  return(list(x_train = x_train, y_train = y_train,
             x_test = x_test, y_test = y_test))
}

```

```

# Lag and split scaled data into train and test
test_size <- 0.2
lookback <- 60

values_split <- lagsplit_data(values, test_size, lookback)

```

```

## x_train.shape = 2875
## x_train.shape = 59
## y_train.shape = 2875
## y_train.shape = 1
## x_test.shape = 719
## x_test.shape = 59
## y_test.shape = 719
## y_test.shape = 1

```

```

x_train <- values_split$x_train
y_train <- values_split$y_train
x_test <- values_split$x_test
y_test <- values_split$y_test

```



# Building the LSTM

```
# Hyperparameters
input_dim <- 1
hidden_dim <- 60 # Units in LSTM (cf. nodes in NN)
# hidden_dim <- round(.5*lookback)
num_layers <- 2
output_dim <- 1
num_epochs <- 7 #50

# Reshape the train and test data into 3D tensor shape
x_train <- array_reshape(x_train, c(dim(x_train)[1], lookback-1, input_dim))
x_test <- array_reshape(x_test, c(dim(x_test)[1], lookback-1, input_dim))

# Define the LSTM model using Keras
model <- keras_model_sequential() %>%
  layer_lstm(units = hidden_dim, return_sequences = TRUE, input_shape = c(lookback-1, input_dim)) %>%
  layer_lstm(units = hidden_dim) %>%
  layer_dense(units = output_dim)

# Compile the model using the mse loss and the Adam optimizer
model %>% compile(loss = "mean_squared_error", optimizer = optimizer_adam(learning_rate = 0.01))
```

**QUESTION :** What does the learning rate do?

# Training the LSTM

```
# Train the model

# Include all data for validation
fitted <- model %>% fit(x_train, y_train, epochs = num_epochs, batch_size = 16, validation_data = list(x_test, y_test))
```

```
## Epoch 1/7
## 180/180 - 7s - 37ms/step - loss: 0.0633 - val_loss: 0.0227
## Epoch 2/7
## 180/180 - 4s - 22ms/step - loss: 0.0261 - val_loss: 0.0223
## Epoch 3/7
## 180/180 - 4s - 24ms/step - loss: 0.0249 - val_loss: 0.0333
## Epoch 4/7
## 180/180 - 3s - 19ms/step - loss: 0.0235 - val_loss: 0.0173
## Epoch 5/7
## 180/180 - 3s - 18ms/step - loss: 0.0227 - val_loss: 0.0183
## Epoch 6/7
## 180/180 - 4s - 21ms/step - loss: 0.0213 - val_loss: 0.0367
## Epoch 7/7
## 180/180 - 4s - 23ms/step - loss: 0.0228 - val_loss: 0.0182
```

```
# Extract predictions from the estimated model
y_train_pred <- model %>% predict(x_train)
```

```
## 90/90 - 1s - 14ms/step
```

```
y_test_pred <- model %>% predict(x_test)
```

```
## 23/23 - 0s - 10ms/step
```

```
# Rescale the predictions and original values
y_train_pred_orig <- y_train_pred * sd(dt$val) + mean(dt$val)
y_train_orig <- y_train * sd(dt$val) + mean(dt$val)
y_test_pred_orig <- y_test_pred * sd(dt$val) + mean(dt$val)
y_test_orig <- y_test * sd(dt$val) + mean(dt$val)
```

## Compute performance metrics

```
# Train MSE and ME
# mse <- fitted$metrics$val_loss[length(fitted$metrics$val_loss)]
rmse <- (mean((y_train_pred_orig - y_train_orig)^2))^0.5
me <- mean(y_train_pred_orig - y_train_orig)
summary(y_train_orig)
```

```
##           V1
## Min.      : 0.00
## 1st Qu.: 6.00
## Median : 9.20
## Mean      :10.25
## 3rd Qu.:14.50
## Max.      :20.60
```

```
# Train Peak error
y_train_orig_p90 <- quantile(y_train_orig, .90)[[1]]
y_train_orig_p90
```

```
## [1] 17.6
```

```
rmse_p90 <- (mean((y_train_orig[y_train_orig >= y_train_orig_p90] - y_train_pred_orig[y_train_orig >= y_train_orig_p90])^2))^.5
me_p90 <- mean( y_train_orig[y_train_orig >= y_train_orig_p90] - y_train_pred_orig[y_train_orig >= y_train_orig_p90])
mpe_p90 <- 100*mean( (y_train_orig[y_train_orig >= y_train_orig_p90] - y_train_pred_orig[y_train_orig >= y_train_orig_p90])/y_train_orig[y_train_orig >= y_train_orig_p90])

# Test MSE and ME
rmse_test <- (mean((y_test_pred_orig - y_test_orig)^2))^.5
me_test <- mean(y_test_pred_orig - y_test_orig)

# Test Peak error
y_test_orig_p90 <- quantile(y_test_orig, .90)[[1]]
y_test_orig_p90
```

```
## [1] 18.2
```

```
rmse_test_p90 <- (mean(( y_test_orig[y_test_orig >= y_test_orig_p90] - y_test_pred_orig[y_test_orig >= y_test_orig_p90])^2))^.5
me_test_p90 <- mean(y_test_orig[y_test_orig >= y_test_orig_p90] - y_test_pred_orig[y_test_orig >= y_test_orig_p90] )
mpe_test_p90 <- 100*mean( (y_test_orig[y_test_orig >= y_test_orig_p90] - y_test_pred_orig[y_test_orig >= y_test_orig_p90])/ y_test_orig[y_test_orig >= y_test_orig_p90])
summary(y_test_orig)
```

```
##          V1
## Min.     : 0.5
## 1st Qu.: 5.6
## Median : 8.1
## Mean    :10.0
## 3rd Qu.:14.7
## Max.    :19.8
```

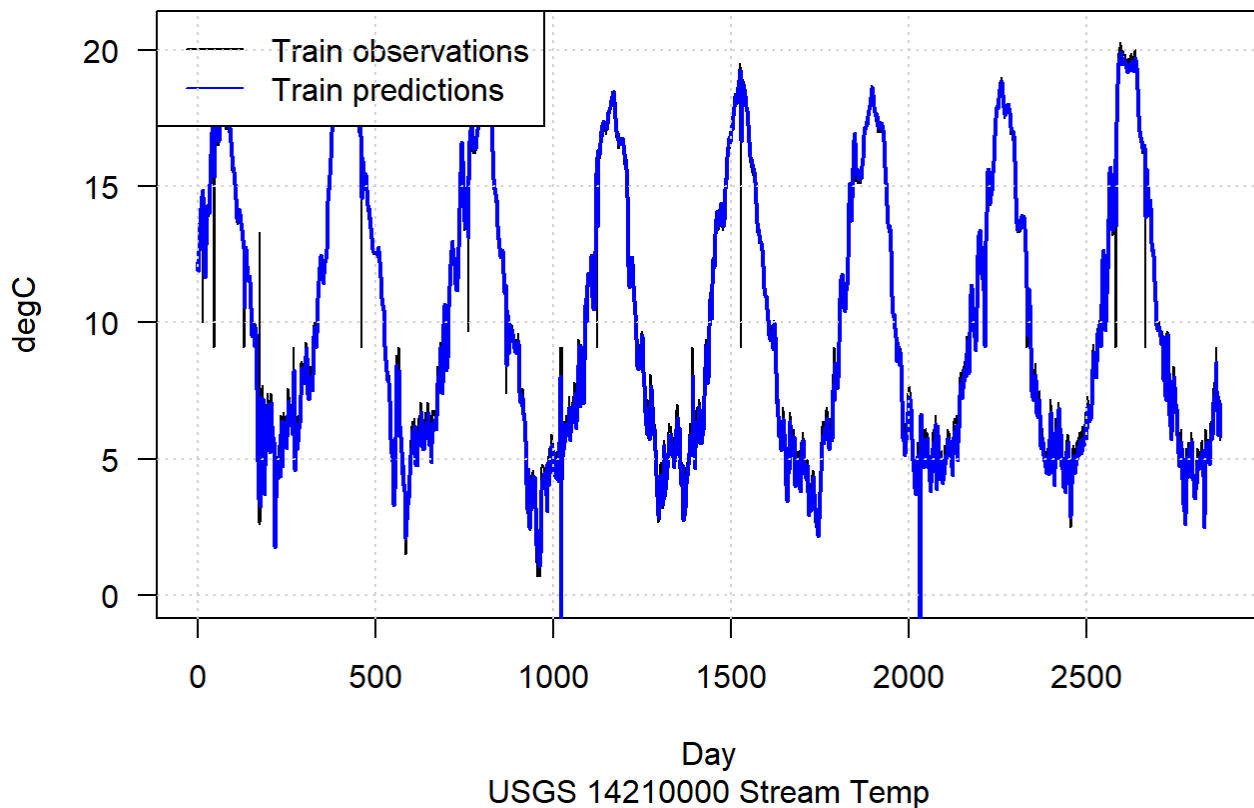
# Plot the LSTM train observations and train predictions

```
pl_title <- "USGS 14210000 Stream Temp"
ylabel <- "degC"

# Set up the plot layout
par(mfrow = c(1, 1))
options(repr.plot.width=15, repr.plot.height=5)

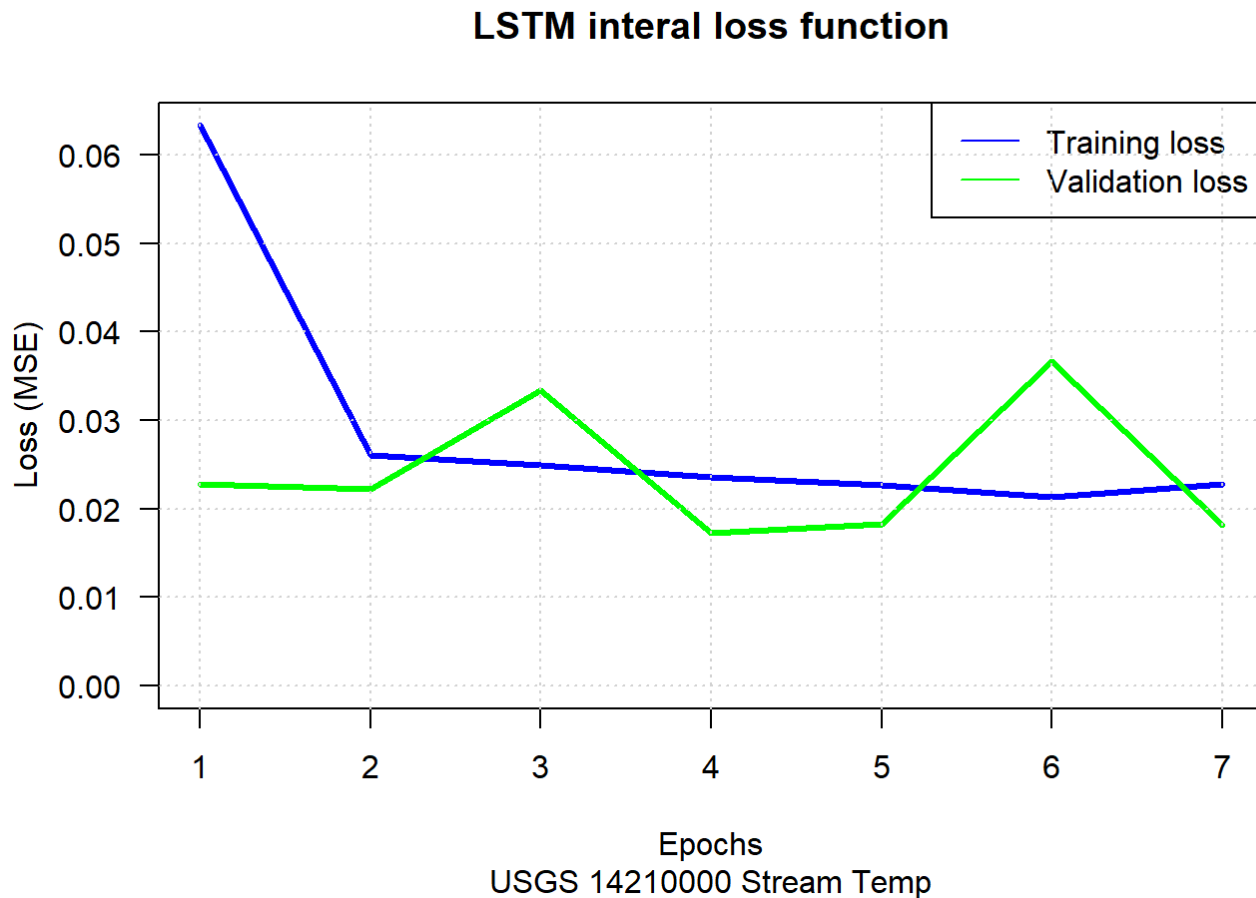
# Plot the train observations and train predictions
plot(y_train_orig, type = "l", main = paste0("LSTM", " Lookback=",lookback," Test size=",t
est_size," Hidden=",hidden_dim," Layers=",num_layers," Epochs=",num_epochs),
      sub = pl_title, col = "black", xlab = "Day", ylab = ylabel, lwd = 1, las=1)
lines(y_train_pred_orig, col = "blue", lwd = 2)
legend(x = "topleft", legend = c("Train observations", "Train predictions"), col = c("blac
k", "blue"), lwd = 1)
grid()
```

**LSTM Lookback=60 Test size=0.2 Hidden=60 Layers=2 Epochs=7**



# LSTM loss functions

```
# Plot the train and test loss
plot(fitted$metrics$loss, type = "l", main = "LSTM interal loss function", xlab = "Epochs", ylab = "Loss (MSE)", sub = pl_title, col = "blue", lwd=3, las=1, ylim = c(0, max(fitted$metrics$loss, fitted$metrics$val_loss)))
lines(fitted$metrics$val_loss, type="l", col = "green", lwd=3, las=1)
legend(x = "topright", legend = c("Training loss", "Validation loss"), col = c("blue", "green"), lwd = 1)
grid()
```

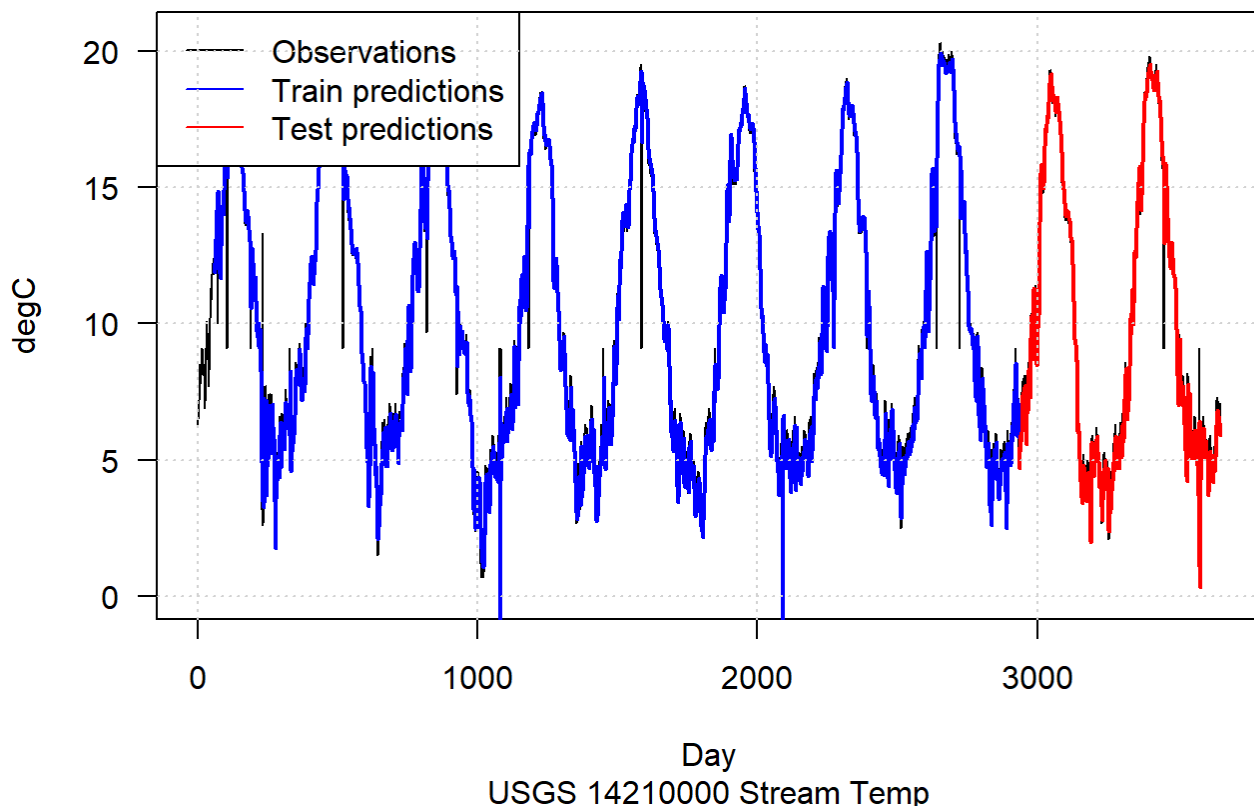


# Plot all observations, LSTM train predictions and testing predictions

```
# Appending the testing predictions to the train predictions
y_train_pred_orig_shifted <- c(rep(NA, lookback), y_train_pred_orig[,1])
shift <- lookback + length(y_train_pred_orig)
y_test_pred_orig_shifted <- c(rep(NA, shift), y_test_pred_orig[,1])
# y_train_pred_orig_shifted <- c(rep(NA, lookback), y_train_pred_orig)
# shift <- lookback + length(y_train_pred_orig)
# y_test_pred_orig_shifted <- c(rep(NA, shift), y_test_pred_orig)

# Plot all observations, train predictions and testing predictions
options(repr.plot.width = 12, repr.plot.height = 8)
plot(dt$val, type = "l", main = paste0("LSTM", " Lookback=",lookback," Test size=",test_si
ze," Hidden=",hidden_dim," Layers=",num_layers," Epochs=",num_epochs, "\nRMSE=", round(rms
e,2), " MPE_p90=", round(mpe_p90,1), "% RMSE_test=", round(rmse_test,2), " MPE_test_p90=",
round(mpe_test_p90,1),"%"),
      sub = pl_title,col = "black", xlab = "Day", ylab = ylabel, lwd = 1, las=1)
lines(y_train_pred_orig_shifted, col = "blue",lwd = 2)
lines(y_test_pred_orig_shifted, col = "red",lwd = 2)
legend(x = "topleft", legend = c("Observations", "Train predictions","Test predictions"),
      col = c("black","blue" ,"red"), lwd = 1)
grid()
```

**LSTM Lookback=60 Test size=0.2 Hidden=60 Layers=2 Epochs=7  
RMSE=0.69 MPE\_p90=0.5% RMSE\_test=0.67 MPE\_test\_p90=0.6%**



Some of the unique features of LSTM neural networks are:

- Long term memory: LSTMs can remember hidden states and register dependencies over long sequences of time
- No vanishing gradient problem: During backpropagation, output is multiplied by the “forget” gate, not by the weights
- Not sensitive to gaps in the data

**QUESTION :** When would you prefer LSTMs for time series forecasting vs a benchmark method?

EOF